

# 1 Basic

## 1.1 Background

Moore's law has given us dramatic increases in hardware system capability, but the *verification gap* in digital logic development process restricts wider, more diverse application. In software, testing costs are proportional to the amount of new code in a design; in hardware, verification costs are proportional to the size of the finished design (including all instances of reused libraries). How can we make the development of *verified* designs more efficient and predictable?

Compilation has few compile checks for invalid concurrent access and proof checks of invariants. While module reuse simplifies design effort, it retains an expensive re-verification burden. Using vendor tooling, the FPGA synthesis process is slow and unstable, preventing developers from lightweight, frequent testing of design changes during the authorship process. No standard debugging when doing on-device testing of interactions with external components. Pipelining of the design must be done manually. Runtime testing (verification) is relied upon as a key aspect of product delivery quality.

All of the above issues are similar to those afflicting software development in the 1960's, leading to what has been described as "the software crisis". Although the underlying model is considerably different from software, analysis of hardware using modern frameworks leads to robust solutions for each of the above issues. All of these issues have a direct impact on both design development cost and on project schedule uncertainty. By improving the agility of the design process, hardware design can move away from the problematic waterfall development process and more easily accomodate evolving application requirements.

In this project, we try to solve many of the major productivity issues in producing validated digital logic designs, enabling wider and more efficient deployment of sophisticated hardware solutions.

## 1.2 Approach

AtomicC is a timed, structural hardware description language for the high level specification of algorithms to be instantiated directly in hardware. AtomicC extends C++ with support for Guarded Atomic Actions [1, 2, 3]: Bluespec-style[4] modules, rules, interfaces, and methods.

Finding errors early – compile time checks:

- check that there are no dynamic concurrency interleaving errors. By conveying developer intent by grouping signal assertions into transactions with guards, we can verify at compile time that there are no invalid combinations at runtime.
- By grouping signal assertions as atomic transactions, we can use proofs to check that user design specific variants are preserved at all times, rather than attempting to check through runtime testing.
- Refinement
- Safety. A safety property asserts that nothing "bad" happens throughout execution
  - Invariant: Basis. Inductive step.
  - Deadlock freedom

- Liveness. A liveness property asserts that something "good" eventually does happen.
  - Termination
  - Response

Example: declarations in Unity:

- Methods: *requires* and *ensures*
- Rules: *Invariant* and *Decrements*
- SystemVerilog: *assume* and *assert. restrict*

Modularize – reuse without reverification:

- contemporary systems contain many components, often designed by distributed teams of engineers. To lower the burden of retesting, these components need to be able to guarantee that runtime operation requirements are preserved.
- By allowing modules to hold off invocations of callers, these module can guarantee that their operating environments are always consistent, hence yield consistent results
- Module reuse without resynthesis: generated Verilog code is reused without regeneration
- Parameterized generated code: propagation of module parameterization from source through to generated Verilog allows reuse without generation for each parameter instance.

Timing closure – simplifying retiming:

- timing closure issues can be separated to 2 groups: interblock and intrablock. By using self-timed interfaces intrablock, long timing closure paths can be eliminated
- simple decoration of source expression trees can be used to generate fifo pipelined implementation of transactions.

Faster edit-compile-test cycles:

- By physically partitioning the design into connected, placed regions, only the logic block that has been actually changed needs to be recompiled. Modular reuse through pre-placed modules

Richer debug environment:

- hardware/software interactions are facilitated by a NOC DMA to processor memory, smoothing interactions between high clock speed/narrow thread software and low clock speed/wide thread hardware.
- runtime support of 'printf' into logs for user logic simplifies the debugging for 'did we get here' type issues.
- distributed logging for signal tracing, supported by simple forwarding of traced signals to software host.

The language is designed for the construction of **modules**[5] that are *correct-by-construction composable*: validated smaller modules can be aggregated to form a larger validated module:

- All state elements in the hardware netlist are explicit in the source code of the design. All module data is private to the module, accessible externally only by method invocation.
- Module interactions are performed with latency insensitive[6, 7] **method** calls, allowing methods to enforce invocation pre-conditions and transitive support for stalling.

- An **interface** is a named collection of method signatures, defining the behavior of an abstract data type(ADT). Modules can declare multiple **interfaces**, giving each interface an explicit name, giving flexibility in coupling with other modules. Interfaces can be exported (defined in the module) or imported (used in the module, but defined externally), giving flexibility in algorithm representation.
- In AtomicC, user operations are written as SSA transactions, called **rules**. Since the compiler can statically analyze the read and write sets of the rule as well as the invocation conditions, it can guarantee that the generated code always executes in a Sequentially Correct (SC) manner: concurrent execution of transactions can be guaranteed to be isolated.
- Like Connectal[8], AtomicC designs may include both hardware and software components, using interfaces to specify type safe communication. The AtomicC compiler generates the code and transactors to pass arguments between hardware and software.

The AtomicC compiler generates a single Verilog module for each defined AtomicC module. Existing Verilog modules can be called from and can call AtomicC generated modules. Standard Verilog backend tools are used to synthesize the resulting ASIC or FPGA. Particular emphasis is put on making the Verilog output both readable and stable to incremental change. Incremental source code changes produce locally incremental generated code changes, easing the management of ECOs and version control on successive releases (for example, management of Verilog output files using git repositories).

Of course, if the underlying algorithm is not designed to allow parallel execution of incremental computations, it will perform poorly and there is nothing the compiler can do to help. AtomicC allows the user to focus solely on the algorithm itself, without the burden of the bare mechanics of orchestrating concurrent consistency, increasing quality and productivity.

### 1.3 Future work

Need to describe multi-cycle rules and pipelining.

Need to describe joining rules

Need to have a way to support sequencing of operations

Need to have a way to support model checking (say 'module B is a behavioral description of module A') Show example with diff eqn solver from Sharp thesis.

Multiple clock domains

Coding FSMs as 'case' statements in Verilog (to integrate with verification tools).

Physical partitioning is used to separate design into separately synthesized pieces, connected using "long distance" signalling. Parallel synthesis; bitstreams combined.

## 2 Modules

The basic building block of AtomicC is the module declaration, made of 3 parts:

- Instantiation of state elements used by the module,
- Interface declarations for interacting with other modules,
- Rules, which group assignment statements and method invocations into atomic transactions.

### 2.1 Module interface definition

An **interface** is a named collection of method signatures, defining the behavior of an abstract data type (ADT).

There are 2 types of methods:

- **Value method functions** allow *inspection* of module state elements.
- **Action method procedures** allow *modification* (including inspection) of state elements, can take parameters and do not have return values. A compiler generated **valid** signal indicates that the caller wishes to perform the method invocation.

Both value and action methods use a compiler generated **ready** signal to indicate when the callee is available and stall scheduling of the calling transaction until execution pre-conditions are satisfied.

[*Example:*

```
__interface EchoRequest {  
    void say(__int(32) v);  
    void say2(__int(16) a, __int(16) b);  
};
```

— *end example*]

### 2.2 Module declaration and definition

Modules are independently compiled. Rule and interface method scheduling logic is generated as part of the generated module. Scheduling constraints (read set, write set and relation to other scheduled elements) are generated into a metadata file, allowing schedule consistency checks between modules to be verified by the linker.

The design is separated into modules that can export and import interfaces to other modules. Each source language module compiles into a single verilog module. Modules are independantly compiled, depending only on the interface definitions for referenced modules. Referencing modules do not depend on the internal implementation of referenced modules, even if they textually exist in the same compilation unit. Scheduling of rules in a module is performed "inside out", with the resulting schedule dependancies written to a metadata file during compilation.

Exported interfaces can be used in several ways:

- invoked directly by the instantiator of the module,

- forwarded transparently, becoming another exported interface of the instantiating module,
- 'connected' to an 'interface reference' of another module in the instantiating scope.

[*Example:*

```
__module Echo {
    EchoRequest    request;           // exported interface (defined by this module)
    EchoIndication *indication;       // imported interface (defined by the instantiator of this module)
    bool busy;
    __int(32) itemSay;
    ...
    // implementation of method request.say(). Note the guard "if (!busy)".
    void request.say(__int(32) v) if(!busy) {
        itemSay = v;
        ...
    }
    void request.saw(__int(16) a, __int(16) b) if(!busy) {
        ...
    }
};
```

— *end example*]

To reference a module from a separate compilation unit, use "`__emodule`". External module definitions need only specify the exported/imported interfaces.

[*Example:*

```
__emodule EchoResponder {
    EchoIndication indication;           // exported interface
};
```

— *end example*]

## 2.3 Guard clauses on module interface methods

- 1 Method definitions in `__module` declarations have the form:

*atomicc-method-definition:*

*decl-specifier-seq<sub>opt</sub> interface-qualifier-seq identifier parameters-and-qualifiers function-body*

*interface-qualifier:*

*identifier .*

*interface-qualifier-seq:*

*interface-qualifier*

*interface-qualifier-seq interface-qualifier*

*atomicc-function-body:*

*if-guard<sub>opt</sub> compound-statement*

*if-guard:*

**if** ( *condition* )

Rules are only ready to fire if the rule's guard is true and all the guards on methods invoked within the rule are also true.

```
void request.say(__int(32) v) if(!busy) {
    itemSay = v;
    ...
}
```

## 2.4 Connecting exported interfaces to imported references

The `__connect` statement allows exported interface declarations to be connected with imported interface references between objects within a module declaration.

*connect-declaration:*

```
__connect identifier = identifier ;
```

[*Example:*

AtomicC example

```
__interface ExampleRequest {
    void say(__int(32) v);
};

__module A {
    ExampleRequest callIn;
};

__module B {
    ExampleRequest *callOut;
};

__module C {
    A consumer;
    B producer;
    __connect producer.callOut = consumer.callIn;
};
```

BSV example

```
BSV example
BSV example
BSV example
BSV example
BSV example
```

— *end example*]

Comparison with BSV:

- The declaration for 'A' is just like BSV. In BSV, the declaration for B requires the interface instance for 'callOut' be passed in as an interface parameter (forcing a textual ordering to the source code declaration sequence).
- In AtomicC, the interfaces are stitched together outside in any convenient sequence in a location where both the concrete instances for A and B are visible.

## 2.5 Exporting interfaces from contained objects

In a design, there are times when the engineer wishes to declare an object locally, but allow external modules to access specific interfaces of the local object. This is done by declaring an interface to the containing object of compatible type and just 'assigning' the local object's interface to it.

[*Example:*

```
__module CWrapper {
    A consumer;
    ExampleRequest request = A.callIn;
};
```

— *end example*]

CWrapper just forwards the interface 'request' down into the instance 'consumer'.

## 2.6 Syntax extension to C++

*atomicc-class-key:*

```
__interface
__emodule
__module
```

## 3 Statements

AtomicC does not attempt to emulate the serialized execution behavior of all C++ constructs in hardware. Instead it uses a subset of the C++ language to specify code blocks that have at most one enabled assignment to any state element. This form is called *static single assignment* form, or SSA form[9]. In the eventual runtime execution, these assignments are all made in a single clock cycle, when the rule or method is enabled.

Constant bound "for" statements that can be fully unrolled are supported.

Since AtomicC does not generate logic to orchestrate sequential execution behavior from language constructs, traditional C++ statements with non-static control flow behavior are not supported.

Examples include:

- Non-constant bound "for" statements.
- "do", "while" statements
- Usages of "goto" that result in a cyclic directed graph of execution blocks
- Method and function calls that are not inlinable at compilation time (for example, recursion is prohibited)

### 3.1 \_\_\_\_rule

Rules specify a group of operations that must execute transactionally: when a rule's guard is satisfied, then it is ready to fire.

Module behavioral statements are encapsulated into transactions (**rules**) following ACID semantics [10, 11]; all rules executed during a given clock cycle are *sequentially consistent* (SC) [12], guaranteeing each rule executes independently of any other rules executing at the same time [11, Sec. 7.1].

*rule-statement:*

`__rule identifier if-guardopt compound-statement`

[*Example:*

```
__rule respond_rule if (responseAvail) {  
    fifo->out.deq();  
    ind->heard(fifo->out.first());  
}
```

— *end example*]

### 3.2 Model details

C block semantics do not correctly process the 2 statements: `a = b; b = a;`. (binding of read values should occur at beginning of block, so that it is clear the 2nd assign refers to the 'previous' value). Thinking again: if we retain C semantics, we have: `temp = a; a = b; b = temp;`, which gives the correct value mapping.

To preserve the standard interpretation of C++ source code in methods and rules, a **modification in-private** [13, Sec. 3.2] execution model is used:

1. Wrap each rule/method with prelude code and postprocessing code
2. Prelude code: For each state element **A** in module, add the declaration of a shadow item:

```
decltype(this->A) A = this->A;    // create shadow of state element
```

3. Postprocessing code: For each state element actually written during execution of the code block:

```
this->A = A;                      // update state elements only at end of code block
```

All assignments in the postprocessing section occur on a single clock cycle.



# 4 External interfacing

## 4.1 Exporting interfaces for use by software

In systems that have both hardware and software components, there is a need to marshal/demarshal parameterized method invocations across a hardware bus or network-on-chip (NOC). AtomicC provides this with my decorating the interface declarations with the keyword `__software`.

The use of the `__software` keyword causes the following to be performed:

- The generation of serialization/deserialization code for both software and hardware side modules to allow the method invocations to be performed in each direction
- The generation of header files allowing compilation of software modules that interface with the hardware
- Integration into a modified Connectal execution framework for the orchestration of requests.

[*Example:*

```
__module Echo {
  __software EchoRequest      request;           // exported interface
  __software EchoIndication   *indication;       // imported interface
  bool busy;
  __int(32) itemSay;
  ...
  // implementation of method request.say(). Note the guard "if (!busy)".
  void request.say(__int(32) v) if(!busy) {
    itemSay = v;
    ...
  }
  void request.saw(__int(16) a, __int(16) b) if(!busy) {
    ...
  }
};
```

— *end example*]

[*Example:*

```
#include "EchoIndication.h" // Header file generated by AtomicC
#include "EchoRequest.h"    // Header file generated by AtomicC

class EchoIndication : public EchoIndicationWrapper
{
public:
  virtual void heard(uint32_t v) {
    // user code for handling indication
  }
  EchoIndication(unsigned int id, PortalTransportFunctions *item, void *param) :
    EchoIndicationWrapper(id, item, param) {}
};

int main(int argc, const char **argv)
{
  EchoIndication echoIndication(IfcNames_EchoIndicationH2S, &transportMux, &param);
  EchoRequestProxy echoRequestProxy(IfcNames_EchoRequestS2H, &transportMux, &param);

  // user code for sending requests
  echoRequestProxy->say(42);
}
```

— *end example*]

## 4.2 Interfacing with legacy Verilog modules

To reference a module in verilog, fields can be declared in `__interface` items.

[*Example:*

```
__interface CNCONNECTNET2 {
    __input  __int(1)    IN1;
    __input  __int(1)    IN2;
    __output __int(1)    OUT1;
    __output __int(1)    OUT2;
};
__emodule CONNECTNET2 {
    CNCONNECTNET2 _;
};
```

— *end example*]

This will allow references/instantiation of an externally defined verilog module CONNECTNET2 that has 2 'input' ports, IN1 and IN2, as well as 2 'output' ports, OUT1 and OUT2.

### 4.2.1 Parameterized modules

Verilog modules that have module instantiation parameters can also be declared/referenced.

[*Example:*

```
__interface Mmcme2MMCME2_ADV {
    __parameter const char * BANDWIDTH;
    __parameter float      CLKFBOUT_MULT_F;
    __input  __uint(1)      CLKFBIN;
    __output __uint(1)      CLKFBOUT;
    __output __uint(1)      CLKFBOUTB;
};
__emodule MMCME2_ADV {
    Mmcme2MMCME2_ADV _;
};
```

— *end example*]

This example can be instantiated as:

[*Example:*

```
__module Test {
    ...
    MMCME2_ADV#(BANDWIDTH="WIDE",CLKFBOUT_MULT_F=1.0) mmcm;
    ...
    Test() {
        __rule initRule {
            mmcm._.CLKFBIN = mmcm._.CLKFBOUT;
        }
    }
}
```

— *end example*]

### 4.2.2 Reference syntax

*atomicc-method-declaration:*

*attribute-specifier-seq<sub>opt</sub> pin-type<sub>opt</sub> decl-specifier-seq<sub>opt</sub> member-declarator-list<sub>opt</sub> ;*

*pin-type:*

```
__input
__output
__inout
__parameter
```

[*Example:*

```
__interface <interfaceName> {  
    __input __uint(1) executeMethod;  
    __input __uint(16) methodArgument;  
    __output __uint(1) methodReady;  
}
```

— *end example*]

For '\_\_\_parameter' items, supported datatypes include: "const char \*", "float", "int".

Factoring of interfaces into sub interfaces is also supported.

### 4.2.3 Clock/reset ports

Note that if interface port pins are declared in a module interface declaration, then CLK and nRST are `__not` automatically declared/instantiated. (Since the user needs the flexibility to not require them when interfacing with legacy code).

Note that this also allows arbitrary signals (like the output of clock generators) to be passed to modules as CLK/nRST signals. (For Atomicc generated modules, please note that the default clock/reset signals for a module will always have these names)

### 4.2.4 Import tooling

There is a tool to automate the creation of AtomicC header files from verilog source files.

[*Example:*

```
atomiccImport -o MMCME2_ADV.h -C MMCME2_ADV -P Mmcme2 zynq.lib  
atomiccImport -o VMMCME2_ADV.h -C MMCME2_ADV -P Mmcme2 MMCME2_ADV.v
```

— *end example*]

# 5 Compilation

AtomicC execution consists of the following phases:

- *compilation*

- Parsing, semantic checks,
- Static elaboration,

During *static elaboration*, constructors are executed for statically declared data elements, allowing allocation of new instances of modules and parameterized rule creation. Any C++ constructs may be used, but the resulting netlist must only contain synthesizable components.

from Newton: Static elaboration does not change the semantics (types, evaluation rules) of the Regiment language, it merely opportunistically pushes evaluation forward into compile time.

We need the ability to create state elements programmatically, improving efficiency of the design creation process.

- Translation to an intermediate representation (IR),
  - Verilog netlist generation from the IR,
- *Linking*: Modules across the project are incrementally compiled and unit tested. As modules are reused in varying contexts, it is necessary to validate if restrictions need to be added for concurrent calls to conflicting methods in the module interface. This checking is done by a *linker*, ensuring that rule/method access across a set of Verilog output is free of inter-module schedule conflicts,
  - *Logic synthesis, physical backend processing*. This is performed with existing backend tool flows
  - *Formal verification* using the Coq Proof Assistant, etc.
  - *Verification*: performed using existing backend tooling,
  - *Hardware execution*.
  - generate Verilog
  - synthesize logic, get size
  - pack blocks with sizes
  - P & R

## 5.1 Linking of groups of modules

To verify that an instantiated group of modules has SC compliant execution characteristics, a linker is used to cross check information from the metadata files for each module.

# Annex A

## Scheduling Algorithm

### A.1 Goals

To guarantee *isolation* in transaction systems, some form of concurrency control is necessary. We can see a classification of concurrency control algorithms below (adapted from [14, Sec. 11.2]):

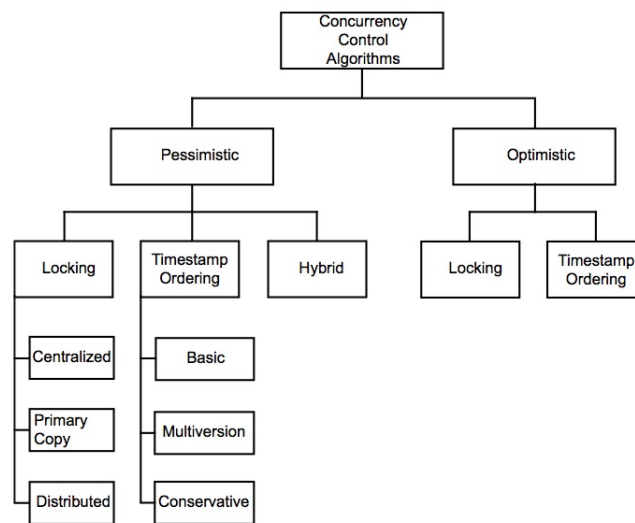


Figure 1 — Concurrency Control Algorithms

ADD TO PICTURE: Static schedule

To guarantee *isolation* in the presence of parallelism in software systems, *dynamic allocation*[11, Sec. 7.3.1] of schedules is used. In hardware design with AtomicC, the set of state elements accessed by a transaction, the operations on these state elements (read-only or write) and the boolean condition when the transaction is performed are all known at compile time. This allows static allocation of **schedules** (sequences of transaction execution) and compile time validation of SC.

**Should we just generate fair schedules in case of conflicts?**

- for write conflicts, use multiway RR arbiter
- for SC conflicts, break cycle with an alternating boolean

### A.2 Signalling

AtomicC uses **valid/ready** *hand-shaking signalling* [15, 16] to invoke action methods, giving both the invoker(master) and invokee(slave) the ability to control invocation execution timing. The master uses the **valid** signal of an action method to show when parameter data is available and the operation should be performed. The method invocation succeeds only when

both **valid** and **ready** are HIGH in the same clock cycle.  
 In TRS notation[1, p. 22]:

$$\pi(M_i) \equiv \text{ready}(M_i) \wedge \text{valid}(M_i).$$

### A.3 Algorithm

The scheduling algorithm is:

- For each module, rules and methods that have some overlap of state element usage (*read set* and *write set*[14, Sec. 10.1.2] [17]) are greedily gathered into *schedule sets*. Since there can be no execution interactions between sets, each set will be independently scheduled.
- A *constraint graph* is a partially-ordered directed graph modeling the schedule sequencing dependencies within a *schedule set*:
  - *nodes* in a constraint graph represent atomic rule and method instances,
  - *edges* represent **write-after-read (WAR)** ordering dependency for a specific storage element[18, Sec. 3].

In addition, each edge has a symbolic boolean *edge condition* for when the the ordering dependency exists: the boolean condition when one rule/method actually reads a given state element and the other actually writes it.
- The transitive closure of these orders on the constraint graph nodes dictate the **schedule** in which each rule must *appear* to execute in order to be considered SC [14, Sec. 11.1]. Of course, since all rules execute in a single cycle, "schedule" does not refer to an actual time sequenced evolution of state, but to a *abstract* "sub-cycle" ordering.
- For each pair of nodes in the constraint digraph, we define the *node condition* between 2 nodes as the conjunction of the *edge conditions* of all the edges between them (i.e., the condition that *any* of the edges causes a dependency). For each cycle in the digraph, we define the *path condition* as the disjunction of the *node conditions* for all sequential pairs of nodes in the cycle (i.e., the condition that *all* the edges, hence the cycle exists).
- Since potential conflicts between methods (called from rules outside the module) and module rules are quite common, if cycle has some method *M* & some rule *R*, then the compiler can automatically rewrite the term *valid(R)* to add a disjunction with the term  $\neg \text{valid}(M)$ , breaking the cycle.
- When the *path condition* is not identically false, a total ordering of the digraph can not be guaranteed and the *schedule set* is not SC. In this case, the compiler or linker reports an error, requiring resolution by the user.

A simple example of a constraint graph is given in A.5, at the end of this document.

Since AtomicC performs scheduling analysis independantly for each declared module, external method invocation conflicts in rules cannot be validated. Schedule processing for external method calls is delayed until the "module group binding" stage of linking, where separately compiled AtomicC output is combined and verified for SC scheduling. Errors and conflicts detected at this stage must be repaired in the module source text and recompiled before proceeding.

## A.4 Previous scheduling work

In Rule Composition[3], scheduling is reformulated in terms of rule composition, leading to a succinct discussion of issues involved, including a concise description of the Esposito and Performance Guarantees schedulers. The resulting schedules are quite close to the user-specified scheduling in AtomicC. In contrast to AtomicC, the Bluespec kernel language they use for analysis also has a sequential composition operator, creating rules that execute for multiple clock cycles.

The Esposito Scheduler[19, 3], is the standard scheduler generation algorithm in the Bluespec Compiler. It uses a heuristic designed to produce a concrete total ordering of rules.

The Performance Guarantees scheduler[20] was proposed to address issues with intra-cycle data passing.

## A.5 Scheduling example

### A.5.1 Source program

```
__interface UserRequest {
    void say(__uint(32) va);
};

__module Order {
    UserRequest request;
    __uint(1) running;
    __uint(32) a, outA, outB, offset;
    void request.say(__uint(32) va) if (!running) {
        a = va;
        offset = 1;
        running = 1;
    }
    __rule A if (!__valid(request.say)) {
        outA = a + offset;
        if (running)
            a = a + 1;
    };
    __rule B if (!__valid(request.say)) {
        outB = a + offset;
        if (!running)
            a = 1;
    };
    __rule C if (!__valid(request.say)) {
        offset = offset + 1;
    };
};
```

### A.5.2 Constraint graph

Sequentially consistent schedules are:

- when 'running == 1': A -> B -> C
- when 'running == 0': B -> A -> C

### A.5.3 Verilog output

```
module Order (input wire CLK, input wire nRST,
    input wire request$say__ENA,
    input wire [31:0]request$say$va,
    output wire request$say__RDY);
    reg [31:0]a, offset, outA, outB;
    reg running;
    assign request$say__RDY = !running;

    always @( posedge CLK) begin
        if (!nRST) begin
            a <= 0;
        end
    end
```

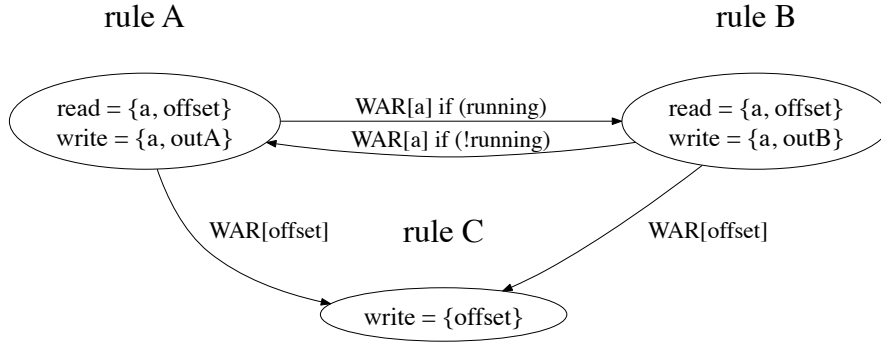


Figure 2 — Simple ordering example

```

offset <= 0;
outA <= 0;
outB <= 0;
running <= 0;
end // nRST
else begin
  if (request$say__ENA == 0) begin // RULE$A__ENA
    outA <= a + offset;
    if (running != 0)
      a <= a + 1;
  end; // End of RULE$A__ENA
  if (request$say__ENA == 0) begin // RULE$B__ENA
    outB <= a + offset;
    if (running == 0)
      a <= 1;
  end; // End of RULE$B__ENA
  if (request$say__ENA == 0) begin // RULE$C__ENA
    offset <= offset + 1;
  end; // End of RULE$C__ENA
  if (request$say__ENA & ( !running )) begin // request$say__ENA
    a <= request$say$va;
    offset <= 1;
    running <= 1;
  end; // End of request$say__ENA
end
end // always @ (posedge CLK)
endmodule

```

#### A.5.4 Intermediate representation output

```

EMODULE l_ainterface_OC_UserRequest {
  METHOD/Action say__ENA ( INTEGER_32 va )
}
MODULE Order {
  INTERFACE l_ainterface_OC_UserRequest request
  FIELD INTEGER_1 running
  FIELD INTEGER_32 a
  FIELD INTEGER_32 outA
  FIELD INTEGER_32 outB
  FIELD INTEGER_32 offset
  METHOD/Rule/Action RULE$A__ENA if (((request$say__ENA) != (0)) ^ (1))) {
    STORE :outA = (a) + (offset)
    STORE ((running) != (0)):a = (a) + (1)
  }
  METHOD/Rule/Action RULE$B__ENA if (((request$say__ENA) != (0)) ^ (1))) {
    STORE :outB = (a) + (offset)
    STORE ((running) != (0)) ^ 1:a = 1
  }
  METHOD/Rule/Action RULE$C__ENA if (((request$say__ENA) != (0)) ^ (1))) {
    STORE :offset = (offset) + (1)
  }
}

```



```
METHOD/Action request$say__ENA ( INTEGER_32 va ) if (((running) != (0)) ^ (1))) {  
    STORE :a = request$say$va  
    STORE :offset = 1  
    STORE :running = 1  
}  
}
```

# Annex B

## Intermediate Representation

### B.1 Module Definitions

*atomicc-module-definition:*

*module-type module-name { module-body-list<sub>opt</sub> }*

*module-type:*

MODULE

EMODULE

INTERFACE

STRUCT

SERIALIZE

*module-body:*

*module-body-definition*

*module-body-definition module-body*

*module-body-definition:*

*field-definition*

*param-definition*

*method-definition*

*interface-definition*

*interface-connect-definition*

*field-definition:*

FIELD *field-option-list<sub>opt</sub>* *field-type* *field-name*

*param-definition:*

PARAMS *field-name* < *param-value* >

*field-option-list:*

*field-option*

*field-option field-option-list*

*field-option:*

/Ptr

/shared

/Count *numeric-constant*

*verilog-field-options*

*verilog-field-options:*

/parameter

/input

/output

/inout

*interface-definition:*

INTERFACE *interface-option-list<sub>opt</sub>* *interface-type* *interface-name*

*interface-connect-definition:*

INTERFACECONNECT *connect-option-list*<sub>opt</sub> *interface-name-l interface-name-r interface-type*

*connect-option-list:*

*connect-option*

*connect-option connect-option-list*

*connect-option-list:*

/Forward

## B.2 Method Definitions

*method-definition:*

METHOD *method-name method-options { method-contents-list }*

*method-options:*

*method-flag-list*<sub>opt</sub> *method-param-list*<sub>opt</sub> *method-return*<sub>opt</sub> *method-guard*<sub>opt</sub>

*method-flag-list:*

*method-flag*

*method-flag method-flag-list*

*method-flag:*

(/Rule)

(/Action)

*param-list:*

*param-definition*

*param-definition , param-list*

*param-definition:*

*param-type param-name*

*method-return:*

*return-type = expression*

*method-guard:*

if *boolean-expression*

*method-contents-list:*

*method-contents*

*method-contents method-contents-list*

*method-contents:*

*alloca-definition*

*let-definition*

*store-definition*

*call-definition*

*alloca-definition:*

ALLOCA *alloca-type alloca-name*

*let-definition:*

LET *let-type action-guard*<sub>opt</sub> : *let-target-name = source-expression*

*store-definition:*

STORE *action-guard*<sub>opt</sub> : *store-target-name = source-expression*

*call-definition:*

CALL *call-option*<sub>opt</sub> *action-guard*<sub>opt</sub> : *call-target-name call-param-list*<sub>opt</sub>

*action-guard:*  
    ( *boolean-expression* )

# Bibliography

- [1] J. C. Hoe, “Operation-Centric Hardware Description and Synthesis,” Ph.D. dissertation, MIT, Cambridge, MA, 2000.
- [2] J. C. Hoe and Arvind, “Operation-Centric Hardware Description and Synthesis,” *IEEE TRANSACTIONS on Computer-Aided Design of Integrated Circuits and Systems*, vol. 23, no. 9, September 2004.
- [3] N. Dave, Arvind, and M. Pellauer, “Scheduling as rule composition,” in *Proceedings of the 5th IEEE/ACM International Conference on Formal Methods and Models for Codesign*, ser. MEMOCODE '07. Washington, DC, USA: IEEE Computer Society, 2007, pp. 51–60.
- [4] Bluespec Inc., <http://www.bluespec.com>.
- [5] D. L. Parnas, “On the criteria to be used in decomposing systems into modules,” *Commun. ACM*, vol. 15, no. 12, pp. 1053–1058, Dec. 1972.
- [6] M. C. Ng, K. E. Fleming, M. Vutukuru, S. Gross, Arvind, and H. Balakrishnan, “Airblue: A system for cross-layer wireless protocol development,” in *Proceedings of the 6th ACM/IEEE Symposium on Architectures for Networking and Communications Systems*, ser. ANCS '10. New York, NY, USA: ACM, 2010, pp. 4:1–4:11.
- [7] M. Abbas and V. Betz, “Latency insensitive design styles for fpgas,” in *28th International Conference on Field Programmable Logic and Applications, FPL 2018, Dublin, Ireland, August 27-31, 2018*, 2018, pp. 360–367.
- [8] M. King, J. Hicks, and J. Ankcorn, “Software-driven hardware development,” in *Proceedings of the 2015 ACM/SIGDA International Symposium on Field-Programmable Gate Arrays*. ACM, 2015, pp. 13–22.
- [9] B. Alpern, M. N. Wegman, and F. K. Zadeck, “Detecting equality of variables in programs,” in *Proceedings of the 15th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, ser. POPL '88. New York, NY, USA: ACM, 1988, pp. 1–11.
- [10] R. S. Nikhil, “Formal specification of bsv’s elaboration and dynamic semantics,” [https://github.com/rsnikhil/Bluespec\\_BSV\\_Formal\\_Semantics](https://github.com/rsnikhil/Bluespec_BSV_Formal_Semantics), 2015.
- [11] J. Gray and A. Reuter, *Transaction Processing: Concepts and Techniques*. Morgan Kaufmann, 1993.
- [12] L. Lamport, “How to make a multiprocessor computer that correctly executes multiprocess programs,” *IEEE Trans. Comput.*, vol. 28, no. 9, pp. 690–691, Sep. 1979.
- [13] A. Prinz and B. Thalheim, “Operational semantics of transactions,” in *IN CRPITS'17: PROCEEDINGS OF THE FOURTEENTH AUSTRALASIAN DATABASE CONFERENCE ON DATABASE TECHNOLOGIES 2003*, 2003, pp. 169–179.

- [14] M. T. Özsu and P. Valduriez, *Principles of Distributed Database Systems, Third Edition*. Springer, 2011.
- [15] C. Fletcher, “Eecs150: Interfaces: “fifo” (a.k.a. ready/valid),” <https://inst.eecs.berkeley.edu/~cs150/Documents/Interfaces.pdf>, 2009.
- [16] L. ARM, “Amba axi and ace protocol specification,” <https://developer.arm.com/docs/ih0022/d/amba-axi-and-ace-protocol-specification-axi3-axi4-and-axi4-lite-ace-and-ace-lite>, 2011.
- [17] D. Rosenkrantz, R. Stearns, and P. Lewis II, “Consistency and serializability in concurrent database systems,” *SIAM Journal on Computing*, vol. 13, no. 3, pp. 508–530, 1984.
- [18] H. W. Cain, M. H. Lipasti, and R. Nair, “Constraint graph analysis of multithreaded programs,” in *Proceedings of the 12th International Conference on Parallel Architectures and Compilation Techniques*, ser. PACT ’03. Washington, DC, USA: IEEE Computer Society, 2003, pp. 4–.
- [19] T. Esposito, M. Lis, R. Nanavati, J. Stoy, and J. Schwartz, “System and method for scheduling TRS rules,” United States Patent US 133051-0001, February 2005.
- [20] D. L. Rosenband and Arvind, “Hardware Synthesis from Guarded Atomic Actions with Performance Specifications,” in *Proceedings of ICCAD’05*, San Jose, CA, 2005.